

STDIN_FILENO

File number of *stdin*. It is 0.

STDOUT_FILENO

File number of *stdout*. It is 1.

STDERR_FILENO

File number of *stderr*. It is 2.

Type Definitions

The `size_t`, `ssize_t`, `uid_t`, `gid_t`, `off_t` and `pid_t` types are defined as described in [<sys/types.h>](#).

The `useconds_t` type is defined as described in [<sys/types.h>](#).

The `intptr_t` type is defined as described in [<inttypes.h>](#).

Declarations

The following are declared as functions and may also be defined as macros. Function prototypes must be provided for use with an ISO C compiler.

```
int      ✓access(const char *, int);
unsigned int alarm(unsigned int);
int      brk(void *);
int      chdir(const char *);
int      chroot(const char *); (LEGACY)
int      ✓chown(const char *, uid_t, gid_t);
✓int     ✓close(int);
size_t   confstr(int, char *, size_t);
char     *crypt(const char *, const char *);
char     *ctermid(char *);
char     *cuserid(char *); (LEGACY)
int      dup(int);
int      dup2(int, int);
void     encrypt(char[64], int);
int      exec1(const char *, const char *, ...);
int      execl(const char *, const char *, ...);
int      execlp(const char *, const char *, ...);
int      execlv(const char *, char *const []);
int      execve(const char *, char *const [], char *const []);
int      execvp(const char *, char *const []);
void     ✓exit(int);
int      ✓fcntl(int, uid_t, gid_t);
int      fcntl(int);
int      fdatasync(int);
pid_t    ✓fork(void);
long int fpathconf(int, int);
int      fsync(int);
int      ftruncate(int, off_t);
char     *getcwd(char *, size_t);
int      getdtablesize(void); (LEGACY)
gid_t    getegid(void);
uid_t    geteuid(void);
gid_t    ✓getgid(void);
int      getgroups(int, gid_t []);
long     gethostid(void);
char     *getlogin(void);
int      getlogin_r(char *, size_t);
int      getopt(int, char *const [], const char *);
int      getpagesize(void); (LEGACY)
char     *getpass(const char *); (LEGACY)
pid_t    getpgid(pid_t);
pid_t    getpgrp(void);
pid_t    ✓getpid(void);
```

unistd.h

fcntl.h

unistd.h

```

✓ pid_t      getpid(void);
✓ pid_t      getsid(pid_t);
uid_t       getuid(void);
char        *getwd(char *);
int         isatty(int);
int         lchown(const char *, uid_t, gid_t);
int         link(const char *, const char *);
int         lockf(int, int, off_t);
off_t       lseek(int, off_t, int);
✓ int        nice(int);
long int    pathconf(const char *, int);
✓ int        pause(void);
int         pipe(int [2]);
ssize_t     pread(int, void *, size_t, off_t);
int         pthread_atfork(void (*)(void), void (*)(void),
                           void (*)(void));
ssize_t     pwrite(int, const void *, size_t, off_t);
✓ ssize_t    read(int, void *, size_t);
int         readlink(const char *, char *, size_t);
int         rmdir(const char *);
void        *sbrk(intptr_t);
✓ int        setgid(gid_t);
✓ int        setpgid(pid_t, pid_t);
pid_t       setpgrp(void);
int         setregid(gid_t, gid_t);
int         setreuid(uid_t, uid_t);
pid_t       setsid(void);
int         setuid(uid_t);
✓ unsigned int sleep(unsigned int);
void        swab(const void *, void *, ssize_t);
int         symlink(const char *, const char *);
void        sync(void);
long int    sysconf(int);
pid_t       tcgetpgrp(int);
int         tcsetpgrp(int, pid_t);
int         truncate(const char *, off_t);
char        *ttyname(int);
int         ttyname_r(int, char *, size_t);
useconds_t  ualarm(useconds_t, useconds_t);
int         unlink(const char *);
int         usleep(useconds_t);
pid_t       vfork(void);
✓ ssize_t    write(int, const void *, size_t);

```

The following external variables are declared:

```

extern char  *optarg;
extern int   optind, opterr, optopt;

```

APPLICATION USAGE

None.

FUTURE DIRECTIONS

None.

SEE ALSO

access(), *alarm()*, *chdir()*, *chown()*, *close()*, *crypt()*, *ctermid()*, *dup()*, *encrypt()*, *environ()*, *exec*,
exit(), *fchdir()*, *fchown()*, *fchmod()*, *fork()*, *fsync()*, *ftruncate()*, *getwd()*, *getgid()*,
getuid(), *getgid()*, *getgroups()*, *gethostid()*, *getlogin()*, *getpgid()*, *getpgrp()*, *getpid()*, *getppid()*,

unistd.h
`void exit(int status)`

- It ends a process and returns a value to its parent.
- status is an integer between 0 to 255. This number is returned to the parent via `wait()` as the exit status of the process.

`exit()` is a library function that calls the system call `_exit()`. `exit()` cleans up the standard I/O streams (`tmpfile()`) before calling `_exit()`.

- Calling `_exit()` instead of `exit()` will bypass this cleanup procedure.
- `wait()` returns the exit status multiplied by 256 (stored in upper 8 bits). The status value is shifted right 8 bit (divided by 256) to obtain the correct value.

`int chown(const char *path, uid_t owner, gid_t group);`

- It sets the owner ID and group ID of the file specified by path to owner & group respectively.
- If owner or group ID is specified as -1 or -1 respectively the corresponding ID of file is unchanged.
- It returns 0 upon successful completion, otherwise -1 is returned.
- It fails if searching permission is denied.

`int fchown(int fd, uid_t, gid_t);`

- `fchown()` has the same effect as `chown()` except that the file path and file description parameter.

pid_t fork();

P ← child's ID

C ← 0

-1 → if Error ⇒ No Resources, ^{insufficient} memory and CHILD_MAX

* pid_t getpid(); Return pid of calling process.

No return value is reserved to indicate an error.

* pid_t getppid(); Return parent process ID of calling process

→ If called ^{No any error} in parent process, it returns shell's process id no.

* int nice(int incr); Return new nice value, -1 in error

→ adds the value of incr to the nice value of the calling process.

→ nice value is a non-negative number.

→ minimum nice value of 0 (zero) are imposed by system.

→ No effect on SCHED_FIFO & SCHED_RR

→ If the process is multi-threaded, the nice value affects all system scope threads in the process.

→ more positive value results in a lower CPU priority

* unsigned int sleep(unsigned int seconds);

- Returns the requested time has elapsed, the value returned will be 0.

- No errors are defined.

- * Pid value will always be unique. Once the process finishes execution its pid value is returned to the kernel which may re-use it for another process sometime later.
- Pid value can't be changed when assigned during execution.
- * `fork()` duplicates the variables which are before `fork()`. It also duplicates even the Global variable (declared before `main()`).
- * `fork()` simply shares the files instead of duplicates.
- * Globally opened file uses the same descriptor. All processes use that one descriptor.

```
int pause(void);
```

SIGCONT

- It suspends the execution of the calling process until it receives a signal.
- There is no successful completion return value. -1 is returned in case of error.
- It fails if a signal is caught by the calling process and control is returned from the signal-catching function.

```
int read(int fd, char* buffer, int nbyte);
```

```
int write(fd, buffer, int nbyte);
```

- fd identifies the I/O channel.
- buffer points to the area in memory where the data is stored for a read() or where the data is taken for a write().
- nbyte defines the maximum number of characters or bytes transferred between the file and the buffer.
- read() write() return the number of bytes transferred.
- There is no limit on nbyte size or depends on SSIZE-MAX. (implementation-dependent).
- A nbyte of 1 is used to transfer a byte at a time so called unbuffered I/O.
- most efficient value for nbyte is the size of the largest physical record the I/O channel is likely to have to handle.

- `read()` & `write()` return a non-negative integer indicating the number of byte actually read or written.
- In case of failure, return `-1` & set `errno` to indicate error.

`int close(int fd);`

- It will deallocate the `fd`. The deallocated `fd` will be available for further use for `open()` & `creat()`. All the outstanding record locks owned by the process on the file associated with `fd` will also be removed.
- It returns `0` on successful completion. Otherwise, it returns `-1` and set `errno` to indicate the error.

`int access(const char* path, int amode);`

- It checks the file pointed by `path` argument is accessible according to the `amode`, using real user ID and real group ID.
- The `amode` constants are defined in `sys/file.h`

<code>F_OK</code>	existence
<code>X_OK</code>	execute
<code>W_OK</code>	write
<code>R_OK</code>	read

* The constant values may be ORed to check more than one access permission.

- It returns `0`, if requested access is permitted. Otherwise returns `-1` to set the `errno`.
- `access()` only answer the question "do I have this permission?"
It can't answer the question "what permission do I have?"

`pid_t setsid(void)`

It creates a new session, if the calling process is not a process group leader. Upon return the calling process will be the session leader of the new session, will be the process group leader of a new process group, and will have no controlling terminal.

The process group ID of the calling process will be set equal to the process ID of the calling process.

- The calling process will be the only process in the new process group and the only process in the new session.
- It returns the value of the process group ID of the calling process on successful completion. Otherwise returns -1 & set `errno` to indicate the error.
- It fails if the calling process is already a process group leader.

`int setpgid(pid_t pid, pid_t pgid)`

- It is used either to join an existing process group or create a new process group within the session of the calling process. The process group ID of a session leader will not change. Upon successful completion, the process group ID of the process with a process ID that matches `pid` will be set to `pgid`.
- As a special case, if `pid` is 0, the process ID of the calling process will be used. Also, if `pgid` is 0, the process group ID of the indicated process will be used.
- Upon successful completion, it returns 0, otherwise -1 is returned and `errno` is set.
- `setpgid(0, 0)` is equivalent to `setpgid(0, 0)`

^{int}
pid_t getpgid(pid_t pid)

- It returns the process group ID of the process whose process ID is equal to pid.
- If pid is equal to 0 (zero), it returns the process group ID of the calling process.
- It fails, it returns or set the errno to indicate the error.
- getpgid() fails if:
 - ① The process whose process ID is equal to pid is not in the same session as the calling process, and the implementation does not allow access to the process group ID of that process from the calling process.
 - ② There is no process with a process ID equal to pid.

Pid_t getpgrp(void)

- It returns the process group ID of the calling process.
- It always successful and no return value is preserved to indicate an error.
- * getpgrp() is equivalent to getpgid(0).
- Each process group is a member of a session and each process is a member of the session of which its process group is a member.

`Pid_t setpgroup(void);`

If the calling process is not already a session leader, `setpgroup()` sets the process group ID of the calling process to the process ID of the calling process. If `setpgroup()` creates a new session, then the new session has no controlling terminal.

- `setpgroup()` has no effect when the calling process is a session leader.
- No any error is defined
- Upon Completion, it returns the process group ID.

`int kill(Pid_t pid, int sig)`

- It will send a signal to a process or a group of processes specified by `pid`. The signal to be sent is specified by `sig` and is either one from the list given in `<signal.h>` or 0.
- If `sig` is 0 (the null signal), error checking is performed but no signal is actual sent. The null signal can be used to check the validity of `pid`.
- If `pid` is greater than 0, `sig` will be sent to the process whose process ID is equal to `pid`.
- If `pid` is 0, `sig` will be sent to all processes whose process group ID is equal to the process group ID of the sender, and for which the process has permission to send a signal.
- If `pid` is -1, `sig` will be sent to all processes for which the process has permission to send that signal.

- If `pid` is negative, but not `-1`, `sig` will be sent to all processes whose process group ID is equal to the absolute value of `pid`, and for which the process has permission to send a signal.
- Upon successful it returns `0` and send `sig` to any of the processes specified by `pid`.
- If `kill()` fails, no signal will be sent, and returns `-1`, and set `errno`.
- `kill()` fails if
 - ① the value of the `sig` is an invalid or unsupported signal number.
 - ② The process does not have permission to send the signal to any receiving process.
 - ③ No process or process group can be found corresponding to that specified by `pid`.

`#include <unistd.h>`

`extern char **environ;`

✓ `int execl(const char *path, const char *arg0, ... /*, (char *)0 */);`
✓ `int execv(const char *path, char *const argv[]);`
✓ `int execl(const char *path, const char *arg0, ... /*, (char *)0, char *const envp[] */);`
✓ `int execve(const char *path, char *const argv[], char *const envp[]);`
✓ `int execlp(const char *file, const char *arg0, ... /*, (char *)0 */);`
✓ `int execvp(const char *file, char *const argv[]);`

- The arguments specified by a program with one of the `exec` functions are passed on to the new process image in the corresponding `main()` arguments.
- The argument `path` points to a pathname that identifies the new process image file.
- The `exec` functions replace the current process image with a new process image. The new image is constructed from a regular, executable file called the new process image file. There is no return from a successful `exec`, because the calling process image is overlaid by the new process image.
- When a C-language program is executed as a result of this call, it is entered as a C-language function call as follows:

```
int main (int argc, char *argv[]);
```

where `argc` is the argument count and `argv` is an array of character pointers to the arguments themselves.

- In addition, the `environ` variable:

```
extern char **environ;
```

- `environ` variable is initialized as a pointer to an array of character pointers to the environment strings.

- ~~Strings~~
- The `argv` and `environ` arrays are each terminated by a null pointer. The null pointer terminating the `argv` array is not counted in `argc`.
 - The argument `file` is used to construct a pathname that identifies the new process image file. If the `file` argument contains a slash character, the `file` argument is used as the pathname for this file. Otherwise, the path prefix for this file is obtained by a search of the directories passed as the environment variable. If this environment variable is not present, the results of the search are implementation-dependent.
 - If the process image file is not a valid executable object, `execlp()` and `execvp()` use the contents of that file as standard input to a command interpreter conforming to `system()`. In this case, the command interpreter becomes the new process image.
 - The arguments represented by `arg0`, ... are pointers to null-terminated character strings. These strings constitute the argument list available to the new process image. The list is terminated by a null pointer. The argument `arg0` should point to a filename that is associated with the process being started by one of the `exec` functions.
 - The argument `argv` is an array of character pointers to null-terminated strings. The last member of this array must be a null pointer. These strings constitute the argument list available to the new process image. The value in `argv[0]` should point to a filename that is associated with the process being started by one of the `exec` functions.
 - The argument `envp` is an array of character pointers to null-terminated strings. These strings constitute the environment for the new process image. The `envp` array is terminated by a null pointer.
 - For those forms not containing an `envp` pointer (eg. `execl`, `execv()`, `execlp()` and `execvp()`), the environment for the new process image is taken from the external variable `environ` in the calling process.
 - The number of bytes available for the new process' combined argument and environment lists is `{ARG_MAX}`. It is implementation-dependent whether null terminators, pointers, and/or any alignment bytes are included in this total.

- nter
- File descriptors open in the calling process image remain open in the new process image, except for those whose close-on-exec flag `FD_CLOEXEC` is set. For those file descriptors that remain open, all attributes of the open file description, including file locks remain unchanged.
 - Directory streams open in the calling process image are closed in the new process image.
 - RETURN VALUE: If one of the exec functions returns to the calling process image, an error has occurred; the return value is -1, and `errno` is set to indicate the error.
 - The new process also inherits at least the following attributes from the calling process image:
 - ✓ nice value (see `nice()`)
 - ✓ `semadj` values (see `semop()`)
 - ✓ process ID
 - ✓ parent process ID
 - ✓ process group ID
 - ✓ session membership
 - ✓ real user ID
 - ✓ real group ID
 - ✓ supplementary group IDs
 - ✓ time left until an alarm clock signal (see `alarm()`)
 - ✓ current working directory
 - ✓ root directory
 - ✓ file mode creation mask (see `umask()`)
 - ✓ file size limit (see `ulimit()`)
 - ✓ process signal mask (see `sigprocmask()`)
 - ✓ pending signal (see `sigpending()`)
 - ✓ `tms_utime`, `tms_stime`, `tms_cutime`, and `tms_cstime` (see `times()`)
 - ✓ resource limits

✓ controlling terminal

✓ interval timers

- All other process attributes defined in this document will be the same in the new and old process images. The inheritance of process attributes not defined by this specification is implementation-dependent.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <process.h>

int main(int argc, char* argv[]) {
    char* arr[] = {
        "ping",
        "google.com",
        NULL
    };

    char* env[] = {
        "TEST=environment variables",
        NULL
    };

    // Without env and PATH variable
    ✓_execv("C:\\\\Windows\\\\System32\\\\ping.exe",
        "C:\\\\Windows\\\\System32\\\\ping.exe", "google.com", NULL);
    ✓_execl("C:\\\\Windows\\\\System32\\\\ping", arr);
    // Without env but with PATH variable
    ✓_execvp("ping", "ping", "google.com", NULL);
    ✓_execlp("ping", arr);
    // With env and PATH variable
    X {
        _execvpe("ping", "ping", "google.com", NULL, env);
        _execlpe("ping", arr, env);
    }
    return 0;
}
```

* `pid_t waitpid(pid_t pid, int *stat_loc, int options)`

→ Arg. `pid_t pid` is greater than 0, It specifies the process ID of a single child process whose status is required.

→ If `pid` is 0, status is required for any child process whose process group ID is equal to that of the calling process.

→ If `pid` is -1, status is requested for any child process. Then `waitpid()` is equivalent to `wait()`.

Options: WCONTINUED (Ask status)

`waitpid()` report the status of any continued child process.

WNOHANG (No more wait)

`waitpid()` shall not suspend the execution of the calling thread if status is not immediately available for child process.

WUNTRACED

child process is stopped but its status is not yet reported. These status reported to the requesting process.

Zombie :- A process that has finished the execution but still has entry in the process table. example - parent is in suspended or sleep state & child exits.

Orphan :- A process whose parent process no more exists eg. Finished or terminated without waiting for its child process to terminate. eg. parent finished but child still executing

wait & waitpid - wait for a child process to stop or terminate.

It forces the parent to suspend execution until the child's finished
x pid_t wait(int *stat_loc)

wait returns the process ID of a child process that finished.

stat_loc is a pointer to an integer where value returned by wait is stored.

⇒ wait() fails by two conditions:

- 1 The process has no children to wait for
- 2 stat_loc points to an invalid address.

⇒ Information returned by wait()

- 1 Ended by exit().
 - 2nd Lowest byte is set to the argument by exit(). Lowest byte is set to 32768
2. Ended by signal.

Lowest byte is set to signal number that ended the process. 2nd Lowest byte is set to 32768