

UNIX System Calls

UNIX System Calls

A system call is just what its name implies -- a request for the operating system to do something on behalf of the user's program. The system calls are functions used in the kernel itself. To the programmer, the system call appears as a normal C function call. However since a system call executes code in the kernel, there must be a mechanism to change the mode of a process from user mode to kernel mode. The C compiler uses a predefined library of functions (the C library) that have the names of the system calls. The library functions typically invoke an instruction that changes the process execution mode to kernel mode and causes the kernel to start executing code for system calls. The instruction that causes the mode change is often referred to as an "operating system trap" which is a software generated interrupt. The library routines execute in user mode, but the system call interface is a special case of an interrupt handler. The library functions pass the kernel a unique number per system call in a machine dependent way -- either as a parameter to the operating system trap, in a particular register, or on the stack -- and the kernel thus determines the specific system call the user is invoking. In handling the operating system trap, the kernel looks up the system call number in a table to find the address of the appropriate kernel routine that is the entry point for the system call and to find the number of parameters the system call expects. The kernel calculates the (user) address of the first parameter to the system call by adding (or subtracting, depending on the direction of stack growth) an offset to the user stack pointer, corresponding to the number of the parameters to the system call. Finally, it copies the user parameters to the "u area" and call the appropriate system call routine. After executing the code for the system call, the kernel determines whether there was an error. If so, it adjusts register locations in the saved user register context, typically setting the "carry" bit for the PS (processor status) register and copying the error number into register 0 location. If there were no errors in the execution of the system call, the kernel clears the "carry" bit in the PS register and copies the appropriate return values from the system call into the locations for registers 0 and 1 in the saved user register context. When the kernel returns from the operating system trap to user mode, it returns to the library instruction after the trap instruction. The library interprets the return values from the kernel and returns a value to the user program.

UNIX system calls are used to manage the file system, control processes, and to provide interprocess communication. The UNIX system interface consists of about 80 system calls (as UNIX evolves this number will increase). The following table lists about 40 of the more important system call:

GENERAL CLASS	SPECIFIC CLASS	SYSTEM CALL
File Structure Related Calls	Creating a Channel	creat() open() close()
	Input/Output	read() write()
	Random Access	lseek()
	Channel Duplication	dup()
	Aliasing and Removing Files	link() unlink()
	File Status	stat() fstat()
	Access Control	access() chmod() chown() umask()
	Device Control	ioctl()
Process Related Calls	Process Creation and Termination	exec() fork() wait() exit()
	Process Owner and Group	getuid() geteuid() getgid() getegid()
	Process Identity	getpid() getppid()
	Process Control	signal() kill() alarm()
	Change Working Directory	chdir()
Interprocess Communication	Pipelines	pipe()
	Messages	msgget() msgsnd() msgrcv() msgctl()
	Semaphores	semget() semop()
	Shared Memory	shmget() shmat() shmdt()

[NOTE: The system call interface is that aspect of UNIX that has changed the most since the inception of the UNIX system. Therefore, when you write a software tool, you should protect that tool by putting system calls in other subroutines within your program and then calling only those subroutines. Should the next version of the UNIX system change the syntax and semantics of the system calls you've used, you need only change your interface routines.]

When a system call discovers an error, it returns -1 and stores the reason the call failed in an external variable named "errno". The "/usr/include/errno.h" file maps these error numbers to manifest constants, and it is these constants that you should use in your programs.

When a system call returns successfully, it returns something other than -1, but it does not clear "errno". "errno" only has meaning directly after a system call that returns an error.

When you use system calls in your programs, you should check the value returned by those system calls. Furthermore, when a system call discovers an error, you should use the "perror()" subroutine to print a diagnostic message on the standard error file that describes why the system call failed. The syntax for "perror()" is:

```
void perror(string)
char string;
```

"perror()" displays the argument string, a colon, and then the error message, as directed by "errno", followed by a newline. The output of "perror()" is displayed on "standard error". Typically, the argument given to "perror()" is the name of the program that incurred the error, argv[0]. However, when using subroutines and system calls on files, the related file name might be passed to "perror()".

There are occasions where you the programmer might wish to maintain more control over the printing of error messages than "perror()" provides -- such as with a formatted screen where the newline printed by "perror()" would destroy the formatting. In this case, you can directly access the same system external (global) variables that "perror()" uses. They are:

```
extern int errno;
extern char *sys_errlist[];
extern int sys_nerr;
```

"errno" has been described above. "sys_errlist" is an array (table) of pointers to the error message strings. Each message string is null terminated and does not contain a newline. "sys_nerr" is the number of messages in the error message table and is the maximum value "errno" can assume. "errno" is used as the index into the table of error messages. Following are two sample programs that display all of the system error messages on standard error.

```

/* errmsg1.c
   print all system error messages using "perror()"
*/

#include <stdio.h>

int main()
{
    int i;
    extern int errno, sys_nerr;

    for (i = 0; i < sys_nerr; ++i)
    {
        fprintf(stderr, "%3d",i);
        errno = i;
        perror(" ");
    }
    exit (0);
}

/* errmsg2.c
   print all system error messages using the global error message table.
*/

#include <stdio.h>

int main()
{
    int i;
    extern int sys_nerr;
    extern char *sys_errlist[];

    fprintf(stderr,"Here are the current %d error messages:\n\n",sys_nerr);
    for (i = 0; i < sys_nerr; ++i)
        fprintf(stderr,"%3d: %s\n", i, sys_errlist[i]);
}

```

Following are some examples in the use of the most often used system calls.

File Structure Related System Calls

The file structure related system calls available in the UNIX system let you create, open, and close files, read and write files, randomly access files, alias and remove files, get information about files, check the accessibility of files, change protections, owner, and group of files, and control devices. These operations either use a character string that defines the absolute or relative path name of a file, or a small integer called a file descriptor that identifies the I/O channel. A channel is a connection between a process and a file that appears to the process as an unformatted stream of bytes. The kernel presents and accepts data from the channel as a process reads and writes that channel. To a process then, all input and output operations are synchronous and unbuffered.

When doing I/O, a process specifies the file descriptor for an I/O channel, a buffer to be filled or emptied, and the maximum size of data to be transferred. An I/O channel may allow input, output, or both. Furthermore, each channel has a read/write pointer. Each I/O operation starts where the last operation finished and advances the pointer by the number of bytes transferred. A process can access a channel's data randomly by changing the read/write pointer.

All input and output operations start by opening a file using either the "creat()" or "open()" system calls. These calls return a file descriptor that identifies the I/O channel. Recall that file descriptors 0, 1, and 2 refer to standard input, standard output, and standard error files respectively, and that file descriptor 0 is a channel to your terminal's keyboard and file descriptors 1 and 2 are channels to your terminal's display screen.

creat()

The prototype for the creat() system call is:

```
int creat(file_name, mode)
char *file_name;
int mode;
```

where file_name is pointer to a null terminated character string that names the file and mode defines the file's access permissions. The mode is usually specified as an octal number such as 0666 that would mean read/write permission for owner, group, and others or the mode may also be entered using manifest constants defined in the "/usr/include/sys/stat.h" file. If the file named by file_name does not exist, the UNIX system creates it with the specified mode permissions. However, if the file does exist, its contents are discarded and the mode value is ignored. The permissions of the existing file are retained. Following is an example of how to use creat():

```

/* creat.c */

#include <stdio.h>
#include <sys/types.h>      /* defines types used by sys/stat.h */
#include <sys/stat.h>      /* defines S_IREAD & S_IWRITE */

int main()
{
    int fd;
    fd = creat("datafile.dat", S_IREAD | S_IWRITE);
    if (fd == -1)
        printf("Error in opening datafile.dat\n");
    else
    {
        printf("datafile.dat opened for read/write access\n");
        printf("datafile.dat is currently empty\n");
    }
    close(fd);
    exit (0);
}

```

The following is a sample of the manifest constants for the mode argument as defined in /usr/include/sys/stat.h:

```

#define S_IRWXU 0000700    /* -rwx----- */
#define S_IREAD 0000400    /* read permission, owner */
#define S_IRUSR S_IREAD
#define S_IWRITE 0000200    /* write permission, owner */
#define S_IWUSR S_IWRITE
#define S_IEXEC 0000100    /* execute/search permission, owner */
#define S_IXUSR S_IEXEC
#define S_IRWXG 0000070    /* ----rwx--- */
#define S_IRGRP 0000040    /* read permission, group */
#define S_IWGRP 0000020    /* write      "      " */
#define S_IXGRP 0000010    /* execute/search "  " */
#define S_IRWXO 0000007    /* -----rwx */
#define S_IROTH 0000004    /* read permission, other */
#define S_IWOTH 0000002    /* write      "      " */
#define S_IXOTH 0000001    /* execute/search "  " */

```

Multiple mode values may be combined by or'ing (using the | operator) the values together as demonstrated in the above sample program.

open()

Next is the open() system call. open() lets you open a file for reading, writing, or reading and writing.

The prototype for the open() system call is:

```
#include <fcntl.h>
```

```
int open(file_name, option_flags [, mode])
char *file_name;
int option_flags, mode;
```

where file_name is a pointer to the character string that names the file, option_flags represent the type of channel, and mode defines the file's access permissions if the file is being created.

The allowable option_flags as defined in "/usr/include/fcntl.h" are:

```
#define O_RDONLY 0      /* Open the file for reading only */
#define O_WRONLY 1     /* Open the file for writing only */
#define O_RDWR  2      /* Open the file for both reading and writing*/
#define O_NDELAY 04     /* Non-blocking I/O */
#define O_APPEND 010    /* append (writes guaranteed at the end) */
#define O_CREAT 00400 /*open with file create (uses third open arg) */
#define O_TRUNC 01000  /* open with truncation */
#define O_EXCL  02000  /* exclusive open */
```

Multiple values are combined using the | operator (i.e. bitwise OR). Note: some combinations are mutually exclusive such as: O_RDONLY | O_WRONLY and will cause open() to fail. If the O_CREAT flag is used, then a mode argument is required. The mode argument may be specified in the same manner as in the creat() system call.

Following is an example of how to use open():

```
/* open.c */

#include <fcntl.h>          /* defines options flags */
#include <sys/types.h>      /* defines types used by sys/stat.h */
#include <sys/stat.h>       /* defines S_IREAD & S_IWRITE */

static char message[] = "Hello, world";

int main()
{
    int fd;
    char buffer[80];

    /* open datafile.dat for read/write access (O_RDWR)
       create datafile.dat if it does not exist (O_CREAT)
       return error if datafile already exists (O_EXCL)
       permit read/write access to file (S_IWRITE | S_IREAD)
    */
    fd = open("datafile.dat", O_RDWR | O_CREAT | O_EXCL, S_IREAD | S_IWRITE);
    if (fd != -1)
    {
        printf("datafile.dat opened for read/write access\n");
        write(fd, message, sizeof(message));
        lseek(fd, 0L, 0); /* go back to the beginning of the file */
        if (read(fd, buffer, sizeof(message)) == sizeof(message))
            printf("\'%s\'" was written to datafile.dat\n", buffer);
        else
            printf("*** error reading datafile.dat ***\n");
        close (fd);
    }
    else
        printf("*** datafile.dat already exists ***\n");
    exit (0);
}
```


close()

To close a channel, use the close() system call. The prototype for the close() system call is:

```
int close(file_descriptor)
int file_descriptor;
```

where file_descriptor identifies a currently open channel. close() fails if file_descriptor does not identify a currently open channel.

read() write()

The read() system call does all input and the write() system call does all output. When used together, they provide all the tools necessary to do input and output sequentially. When used with the lseek() system call, they provide all the tools necessary to do input and output randomly.

Both read() and write() take three arguments. Their prototypes are:

```
int read(file_descriptor, buffer_pointer, transfer_size)
int file_descriptor;
char *buffer_pointer;
unsigned transfer_size;
```

```
int write(file_descriptor, buffer_pointer, transfer_size)
int file_descriptor;
char *buffer_pointer;
unsigned transfer_size;
```

where file_descriptor identifies the I/O channel, buffer_pointer points to the area in memory where the data is stored for a read() or where the data is taken for a write(), and transfer_size defines the maximum number of characters transferred between the file and the buffer. read() and write() return the number of bytes transferred.

There is no limit on transfer_size, but you must make sure it's safe to copy transfer_size bytes to or from the memory pointed to by buffer_pointer. A transfer_size of 1 is used to transfer a byte at a time for so-called "unbuffered" input/output. The most efficient value for transfer_size is the size of the largest physical record the I/O channel is likely to have to handle. Therefore, 1K bytes -- the disk block size -- is the most efficient general-purpose buffer size for a standard file. However, if you are writing to a terminal, the transfer is best handled in lines ending with a newline.

For an example using read() and write(), see the above example of open().

`lseek()`

The UNIX system file system treats an ordinary file as a sequence of bytes. No internal structure is imposed on a file by the operating system. Generally, a file is read or written sequentially -- that is, from beginning to the end of the file. Sometimes sequential reading and writing is not appropriate. It may be inefficient, for instance, to read an entire file just to move to the end of the file to add characters. Fortunately, the UNIX system lets you read and write anywhere in the file. Known as "random access", this capability is made possible with the `lseek()` system call. During file I/O, the UNIX system uses a long integer, also called a File Pointer, to keep track of the next byte to read or write. This long integer represents the number of bytes from the beginning of the file to that next character. Random access I/O is achieved by changing the value of this file pointer using the `lseek()` system call.

The prototype for `lseek()` is:

```
long lseek(file_descriptor, offset, whence)
int file_descriptor;
long offset;
int whence;
```

where `file_descriptor` identifies the I/O channel and `offset` and `whence` work together to describe how to change the file pointer according to the following table:

whence	new position

0	offset bytes into the file
1	current position in the file plus offset
2	current end-of-file position plus offset

If successful, `lseek()` returns a long integer that defines the new file pointer value measured in bytes from the beginning of the file. If unsuccessful, the file position does not change.

Certain devices are incapable of seeking, namely terminals and the character interface to a tape drive. `lseek()` does not change the file pointer to these devices.

Following is an example using lseek():

```
/* lseek.c */

#include <stdio.h>
#include <fcntl.h>

int main()
{
    int fd;
    long position;

    fd = open("datafile.dat", O_RDONLY);
    if ( fd != -1)
    {
        position = lseek(fd, 0L, 2); /* seek 0 bytes from end-of-file */
        if (position != -1)
            printf("The length of datafile.dat is %ld bytes.\n", position);
        else
            perror("lseek error");
    }
    else
        printf("can't open datafile.dat\n");
    close(fd);
}
```

Many UNIX systems have defined manifest constants for use as the "whence" argument of lseek(). The definitions can be found in the "file.h" and/or "unistd.h" include files. For example, the University of Maryland's HP-9000 UNIX system has the following definitions:

from file.h we have:

```
#define L_SET          0      /* absolute offset */
#define L_INCR         1      /* relative to current offset */
#define L_XTND         2      /* relative to end of file */
```

and from unistd.h we have:

```
#define SEEK_SET       0      /* Set file pointer to "offset" */
#define SEEK_CUR       1      /* Set file pointer to current plus "offset" */
#define SEEK_END       2      /* Set file pointer to EOF plus "offset" */
```

The definitions from unistd.h are the most "portable" across UNIX and MS-DOS C compilers.

dup()

The dup() system call duplicates an open file descriptor and returns the new file descriptor. The new file descriptor has the following properties in common with the original file descriptor:

- refers to the same open file or pipe.

- has the same file pointer -- that is, both file descriptors share one file pointer.

- has the same access mode, whether read, write, or read and write.

The prototype for dup() is:

```
int dup(file_descriptor)
int file_descriptor;
```

where file_descriptor is the file descriptor describing the original I/O channel returned by creat(), open(), pipe(), or dup() system calls. dup() is guaranteed to return a file descriptor with the lowest integer value available. It is because of this feature of returning the lowest unused file descriptor available that processes accomplish I/O redirection. The following example shows standard output redirected to a file through the use of the dup() system call:

```

/* dup.c
   demonstrate redirection of standard output to a file.
*/

#include <stdio.h>
#include <fcntl.h>
#include <sys/types.h>
#include <sys/stat.h>

int main()
{
    int fd;

    fd = open("foo.bar", O_WRONLY | O_CREAT, S_IREAD | S_IWRITE );
    if (fd == -1)
    {
        perror("foo.bar");
        exit (1);
    }
    close(1);          /* close standard output */
    dup(fd);           /* fd will be duplicated into standard out's slot */
    close(fd);         /* close the extra slot */
    printf("Hello, world!\n"); /* should go to file foo.bar */
    exit (0);          /* exit() will close the files */
}

```

link()

The UNIX system file structure allows more than one named reference to a given file, a feature called "aliasing". Making an alias to a file means that the file has more than one name, but all names of the file refer to the same data. Since all names refer to the same data, changing the contents of one file changes the contents of all aliases to that file. Aliasing a file in the UNIX system amounts to the system creating a new directory entry that contains the alias file name and then copying the i-number of an existing file to the i-number position of this new directory entry. This action is accomplished by the link() system call. The link() system call links an existing file to a new file.

The prototype for link() is:

```
int link(original_name, alias_name)
char *original_name, *alias_name;
```

where both original_name and alias_name are character strings that name the existing and new files respectively. link() will fail and no link will be created if any of the following conditions holds:

- a path name component is not a directory.
- a path name component does not exist.
- a path name component is off-limits.
- original_name does not exist.
- alias_name does exist.
- original_name is a directory and you are not the superuser.
- a link is attempted across file systems.
- the destination directory for alias_name is not writable.
- the destination directory is on a mounted read-only file system.

Following is a short example:

```
/* link.c
*/

#include <stdio.h>

int main()
{
    if ((link("foo.old", "foo.new")) == -1)
    {
        perror(" ");
        exit(1);          /* return a non-zero exit code on error */
    }
    exit(0);
}
```

unlink()

The opposite of the link() system call is the unlink() system call. unlink() removes a file by zeroing the i-number part of the file's directory entry, reducing the link count field in the file's inode by 1, and releasing the data blocks and the inode if the link count field becomes zero. unlink() is the only system call for removing a file in the UNIX system.

The prototype for unlink() is:

```
int unlink(file_name)
char *file_name;
```

where file_name names the file to be unlinked. unlink() fails if any of the following conditions holds:

- a path name component is not a directory.
- a path name component does not exist.
- a path name component is off-limits.
- file_name does not exist.
- file_name is a directory and you are not the superuser.
- the directory for the file named by file_name is not writable.
- the directory is contained in a file system mounted read-only.

It is important to understand that a file's contents and its inode are not discarded until all processes close the unlinked file.

Following is a short example:

```
/* unlink.c
*/

#include <stdio.h>

int main()
{
    if ((unlink("foo.bar")) == -1)
    {
        perror(" ");
        exit (1);          /* return a non-zero exit code on error */
    }
    exit (0);
}
```

Process Related System Calls

exec

The UNIX system provides several system calls to create and end program, to send and receive software interrupts, to allocate memory, and to do other useful jobs for a process. Four system calls are provided for creating a process, ending a process, and waiting for a process to complete. These system calls are `fork()`, the "exec" family, `wait()`, and `exit()`.

The UNIX system calls that transform an executable binary file into a process are the "exec" family of system calls. The prototypes for these calls are:

```
int execl(file_name, arg0 [, arg1, ..., argn], NULL)
char *file_name, *arg0, *arg1, ..., *argn;
```

```
int execv(file_name, argv)
char *file_name, *argv[];
```

```
int execlp(file_name, arg0 [, arg1, ..., argn], NULL, envp)
char *file_name, *arg0, *arg1, ..., *argn, *envp[];
```

```
int execve(file_name, argv, envp)
char *file_name, *argv[], *envp[];
```

```
int execlp(file_name, arg0 [, arg1, ..., argn], NULL)
char *file_name, *arg0, *arg1, ..., *argn;
```

```
int execvp(file_name, argv)
char *file_name, *argv[];
```

where `file_name` names the executable binary file to be transformed into a process, `arg0` through `argn` and `argv` define the arguments to be passed to the process, and `envp` defines the environment, also to be passed to the process. By convention, `arg0` and `argv[0]` name the last path name component of the executable binary file named by `file_name`. For `execl()`, `execv()`, `execlp()`, and `execve()`, `file_name` must be the fully qualified path name of the executable binary file. However for `execlp()` and `execvp()`, the `PATH` variable is used to find the executable binary file. When the environment is not explicitly given as an argument to an exec system call, the environment of the current process is used. Furthermore, the last array element of both `argv` and `envp` must be null to signify the end of the array.

Unlike the other system calls and subroutines, a successful exec system call does not return. Instead, control is given to the executable binary file named as the first argument. When that file is made into a process, that process replaces the process that executed the exec system call -- a new process is not created. If an exec call should fail, it will return a -1.

Letters added to the end of exec indicate the type of arguments:

- l argn is specified as a list of arguments.
- v argv is specified as a vector (array of character pointers).
- e environment is specified as an array of character pointers.
- p user's PATH is searched for command, and command can be a shell program

Following is a brief description of the six routines that make up the collective family of exec routines:

execl Takes the path name of an executable program (binary file) as its first argument. The rest of the arguments are a list of command line arguments to the new program (argv[]). The list is terminated with a null pointer:

```
execl("/bin/cat", "cat", "f1", "f2", (char *) 0);
execl("a.out", "a.out", (char *) 0);
```

Note that, by convention, the argument listed after the program is the name of the command being executed (argv[0]).

execle Same as execl(), except that the end of the argument list is followed by a pointer to a null-terminated list of character pointers that is passed as the environment of the new program (i.e., the place that getenv() searches for exported shell variables):

```
static char *env[] = {
    "TERM=vt100",
    "PATH=/bin:/usr/bin",
    (char *) 0 };

execle("/bin/cat", "cat", "f1", "f2", (char *) 0, env);
```

execv Takes the path name of an executable program (binary file) as its first argument. The second argument is a pointer to a list of character pointers (like argv[]) that is passed as command line arguments to the new program:

```
static char *args[] = {
    "cat",
    "f1",
    "f2",
    (char *) 0 };

execv("/bin/cat", args);
```

`execve` Same as `execv()`, except that a third argument is given as a pointer to a list of character pointers (like `argv[]`) that is passed as the environment of the new program:

```
static char *env[] = {
    "TERM=vt100",
    "PATH=/bin:/usr/bin",
    (char *) 0 };

static char *args[] = {
    "cat",
    "f1",
    "f2",
    (char *) 0 };

execve("/bin/cat", args, env);
```

`execlp` Same as `execl()`, except that the program name doesn't have to be a full path name, and it can be a shell program instead of an executable module:

```
execlp("ls", "ls", "-l", "/usr", (char *) 0);

execlp() searches the PATH environment variable to find the
specified program.
```

`execvp` Same as `execv()`, except that the program name doesn't have to be a full path name, and it can be a shell program instead of an executable module:

```
static char *args[] = {
    "cat",
    "f1",
    "f2",
    (char *) 0 };

execvp("cat", args);
```

When transforming an executable binary file into a process, the UNIX system preserves some characteristics of the replaced process. Among the items saved by the `exec` system call are:

- The "nice" value for scheduling.
- The process ID and the parent process ID.
- The time left until an alarm clock signal.
- The current working directory and the root directory.
- The file creation mask as established with `umask()`.
- All open files.

The last of these is the most interesting because the shell uses this feature to handle input/output redirection.

fork()

The exec family of system calls transforms an executable binary file into a process that overlays the process that made the exec system call. The UNIX system does not create a new process in response to an exec system call. To create a new process, you must use the fork() system call. The prototype for the fork() system call is:

```
int fork()
```

fork() causes the UNIX system to create a new process, called the "child process", with a new process ID. The contents of the child process are identical to the contents of the parent process.

The new process inherits several characteristics of the old process. Among the characteristics inherited are:

- The environment.
- All signal settings.
- The set user ID and set group ID status.
- The time left until an alarm clock signal.
- The current working directory and the root directory.
- The file creation mask as established with umask().

The child process begins executing and the parent process continues executing at the return from the fork() system call. This is difficult to understand at first because you only call fork() once, yet it returns twice -- once per process. To differentiate which process is which, fork() returns zero in the child process and non-zero (the child's process ID) in the parent process.

exec routines are usually called after a call to fork(). This combination, known as a fork/exec, allows a process to create a child to execute a command, so that the parent doesn't destroy itself through an exec. Most command interpreters (e.g. the shell) on UNIX use fork and exec.

wait()

You can control the execution of child processes by calling wait() in the parent. wait() forces the parent to suspend execution until the child is finished. wait() returns the process ID of a child process that finished. If the child finishes before the parent gets around to calling wait(), then when wait() is called by the parent, it will return immediately with the child's process ID. (It is possible to have more than one child process by simply calling fork() more than once.). The prototype for the wait() system call is:

```
int wait(status)
int *status;
```

where status is a pointer to an integer where the UNIX system stores the value returned by the child process. wait() returns the process ID of the process that ended. wait() fails if any of the following conditions hold:

- The process has no children to wait for.
- status points to an invalid address.

The format of the information returned by wait() is as follows:

If the process ended by calling the exit() system call, the second lowest byte of status is set to the argument given to exit() and the lowest byte of status is set to zeroes.

If the process ended because of a signal, the second lowest byte of status is set to zeroes and the lowest byte of status contains the signal number that ended the process. If the seventh bit of the lowest byte of status is set (i.e. status & 0200 == 0200) then the UNIX system produced a core dump of the process.

exit()

The exit() system call ends a process and returns a value to its parent. The prototype for the exit() system call is:

```
void exit(status)
int status;
```

where status is an integer between 0 and 255. This number is returned to the parent via wait() as the exit status of the process. By convention, when a process exits with a status of zero that means it didn't encounter any problems; when a process exits with a non-zero status that means it did have problems.

exit() is actually not a system routine; it is a library routine that calls the system routine _exit(). exit() cleans up the standard I/O streams before calling _exit(), so any output that has been buffered but not yet actually written out is flushed. Calling _exit() instead of exit() will bypass this cleanup procedure. exit() does not return.

Following are some example programs that demonstrate the use of fork(), exec(), wait(), and exit():

```
/* status.c
   demonstrates exit() returning a status to wait().
*/

int main()
{
    unsigned int status;

    if ( fork () == 0 ) {          /* == 0 means in child */
        scanf ("%d", &status);
        exit (status);
    }
    else {                         /* != 0 means in parent */
        wait (&status);
        printf("child exit status = %d\n", status > 8);
    }
}
```

Note: since wait() returns the exit status multiplied by 256 (contained in the upper 8 bits), the status value is shifted right 8 bits (divided by 256) to obtain the correct value.

```

/* myshell.c
   This program is a simple command interpreter that uses execlp() to
   execute commands typed in by the user.
*/
#include <stdio.h>
#define EVER ;;

int main()
{
    int process;
    char line[81];

    for (EVER)
    {
        fprintf(stderr, "cmd: ");
        if ( gets (line) == (char *) NULL)      /* blank line input */
            exit (0);

        /* create a new process */

        process = fork ();

        if (process > 0)          /* parent */
            wait ((int *) 0);    /* null pointer - return value not saved */
        else if (process == 0)   /* child */
        {
            /* execute program */
            execlp (line, line, (char *) NULL);
            /* some problem if exec returns */
            fprintf (stderr, "Can't execute %s\n", line);
            exit (1);
        }
        else if ( process == -1) /* can't create a new process */
        {
            fprintf (stderr, "Can't fork!\n");
            exit (2);
        }
    }
}

```

The following program demonstrates a practical use of `fork()` and `exec()` to create a new directory. Only the superuser has the permission to use the `mknod()` system call to create a new directory -- an ordinary user cannot use `mknod()` to create a directory. So, we use `fork/exec` to call upon the UNIX system's `mkdir` command that anyone can use to create a directory.

```
/* newdir.c
   create a new directory, called newdir, using fork() and exec().
*/
#include <stdio.h>

int main()
{
    int fd;

    if ( fork() != 0)
        wait ((int *) 0);
    else
    {
        execl ("/bin/mkdir", "mkdir", "newdir", (char *) NULL);
        fprintf (stderr, "exec failed!\n");
        exit (1);
    }

    /* now use newdir */
    if ( (fd = open("newdir/foo.bar", O_RDWR | O_CREAT, 0644)) == -1)
    {
        fprintf (stderr, "open failed!\n");
        exit (2);
    }
    write (fd, "Hello, world\n", 14);
    close (fd);
    exit (0);
}
```

Software Interrupts

signal()

The UNIX system provides a facility for sending and receiving software interrupts, also called SIGNALS. Signals are sent to a process when a predefined condition happens. The number of signals available is system dependent. For example, the University's HP-9000 has 31 signals defined. The signal name is defined in /usr/include/sys/signal.h as a manifest constant.

Programs can respond to signals three different ways. These are:

1. Ignore the signal. This means that the program will never be informed of the signal no matter how many times it occurs. The only exception to this is the SIGKILL signal which can neither be ignored nor caught.
2. A signal can be set to its default state, which means that the process will be ended when it receives that signal. In addition, if the process receives any of SIGQUIT, SIGILL, SIGIOT, SIGEMT, SIGFPE, SIGBUS, SIGSEGV, or SIGSYS, the UNIX system will produce a core image (core dump), if possible, in the directory where the process was executing when it received the program-ending signal.
3. Catch the signal. When the signal occurs, the UNIX system will transfer control to a previously defined subroutine where it can respond to the signal as is appropriate for the program.

You define how you want to respond to a signal with the signal() system call. The prototype is:

```
#include <sys/signal.h>

int (* signal ( signal_name, function ))
int signal_name;
int (* function)();
```

where signal_name is the name of the signal from signal.h and function is any of SIG_IGN, meaning that you wish to ignore the signal when it occurs; SIG_DFL, meaning that you wish the UNIX system to take the default action when your program receives the signal; or a pointer to a function that returns an integer. The function is given control when your program receives the signal, and the signal number is passed as an argument. signal() returns the previous value of function, and signal() fails if any of the following conditions hold:

signal_name is an illegal name or SIGKILL.

function points to an invalid memory address.

Once a signal is caught, the UNIX system resets it to its initial state (the default condition). In general, if you intend for your program to be able to catch a signal repeatedly, you need to re-arm the signal handling mechanism. You must do this as soon after receipt of the signal as possible, namely just after entering the signal handling routine.

You should use signals in your programs to isolate critical sections from interruption.

The state of all signals is preserved across a `fork()` system call, but all caught signals are set to `SIG_DFL` across an `exec` system call.

`kill()`

The UNIX system sends a signal to a process when something happens, such as typing the interrupt key on a terminal, or attempting to execute an illegal instruction. Signals are also sent to a process with the `kill()` system call. Its prototype is:

```
int kill (process_id, signal_name )
int process_id, signal_name;
```

where `process_id` is the ID of the process to be signaled and `signal_name` is the signal to be sent to that process. If `process_id` has a positive value, that value is assumed to be the process ID of the process to whom `signal_name` signal is to be sent. If `process_id` has the value 0, then `signal_name` signal is sent to all processes in the sending process' process group, that is all processes that have been started from the same terminal. If `process_id` has the value -1 and the process executing the `kill()` system call is the superuser, then `signal_name` is sent to all processes excluding process 0 and process 1 that have the same user ID as the process executing the `kill()`. `kill()` fails if any of the following conditions hold:

- `signal_name` is not a valid signal.

- there is not a process in the system with process ID `process_id`.

- even though the process named by `process_id` is in the system, you cannot send it a signal because your effective user ID does not match either the real or effective user ID of `process_id`.

alarm()

Every process has an alarm clock stored in its system-data segment. When the alarm goes off, signal SIGALRM is sent to the calling process. A child inherits its parent's alarm clock value, but the actual clock isn't shared. The alarm clock remains set across an exec.

The prototype for alarm() is:

```
unsigned int alarm(seconds)
unsigned int seconds;
```

where seconds defines the time after which the UNIX system sends the SIGALRM signal to the calling process. Each successive call to alarm() nullifies the previous call, and alarm() returns the number of seconds until that alarm would have gone off. If seconds has the value 0, the alarm is canceled. alarm() has no error conditions.

The following is an example program that demonstrates the use of the signal() and alarm() system calls:

```
/* timesup.c */

#include <stdio.h>
#include <sys/signal.h>

#define EVER ;;

void main();
int times_up();

void main()
{
    signal (SIGALRM, times_up);    /* go to the times_up function */
    alarm (10);                   /* when the alarm goes off. */
    /* set the alarm for 10 seconds */

    for (EVER)                   /* endless loop. */
        ;                       /* hope the alarm works. */
}

int times_up(sig)
int sig;                        /* value of signal */
{
    printf("Caught signal #< %d >\n", sig);
    printf("Time's up! I'm outta here!!\n");
    exit(sig);                   /* return the signal number */
}
```

Interprocess Communication

UNIX System V allows processes to communicate with one another using pipes, messages, semaphores, and shared memory. This section describes how to communicate using pipes.

One way to communicate between two processes is to create a pipeline with the `pipe()` system call. `pipe()` builds the channel, but it is up to you to connect the standard input of one process to the standard output of the other process.

The prototype for `pipe()` is:

```
int pipe (file_descriptors)
int file_descriptors[2];
```

where `file_descriptors[2]` is an array that `pipe()` fills with a file descriptor opened for reading, `file_descriptor[0]`, and a file descriptor opened for writing, `file_descriptor[1]`. `pipe()` fails for the following condition:

there are too many open I/O channels.

Some I/O system calls act differently on pipe file descriptors from the way they do on ordinary files, and some do nothing at all. Following is a summary of these actions:

- write Data written to a pipe is sequenced in order of arrival. Normally, if the pipe becomes full, `write()` will block until enough old data is removed by `read()`. There are no partial writes; the entire `write()` will be completed. The capacity of a pipe varies with the UNIX implementation, but it is always at least 4096 bytes (4K). If `fcntl()` is called to set the `O_NDELAY` flag, `write()` will not block on a full pipe and will return a count of 0. The only way to put an end-of-file on a pipe is to close the writing file descriptor.
- read Data is read from a pipe in order of arrival, just as it was written. Once read, data can't be reread or put back. Normally, if the pipe is empty, `read` will block until at least one byte of data is available, unless the writing file descriptor is closed, in which case the `read` will return a 0 count (the usual end-of-file indication). But the byte count given as the third argument to `read` will not necessarily be satisfied - only as many bytes as are present at that instant will be read, and an appropriate count will be returned. The byte count will never be exceeded, of course; unread bytes will remain for the next `read()`. If the `O_NDELAY` flag is set, a `read()` on an empty pipe will return with a 0 count. This suffers from the same ambiguity as reads on communication lines. A 0 count also means end-of-file.

close Means more on a pipe than it does on a file. Not only does it free up the file descriptor for reuse, but when the writing file descriptor is closed it acts as an end-of-file for the reader. If the read end file descriptor is closed, a write() on the other file descriptor will cause an error. A fatal signal is also normally generated (SIGPIPE - #13).

fcntl This system call sets or clears the O_NDELAY flag, whose effect is described under write and read above.

fstat Not very useful on pipes. The size returned is the number of bytes in the pipe, but this fact is seldom useful. A pipe may be distinguished by a link count of 0, since a pipe is the only source of a file descriptor associated with something not linked into a directory. This distinction might be useful to I/O routines that want to treat pipes specially.

open Not used with pipes.

creat Not used with pipes.

lseek Not used with pipes. This means that if a pipe contains a sequence of messages, it isn't possible to look through them for the message to read next. Like toothpaste in a tube, you have to get it out to examine it, and then there is no way to put it back.

Pipes use the buffer cache just as ordinary files do. Therefore, the benefits of writing and reading pipes in units of a block (usually 512 or 1024 bytes) are just as great. A single write() execution is atomic, so if 512 bytes are written with a single system call, the corresponding read() will return with 512 bytes (if it requests that many). It will not return with less than the full block. However, if the writer is not writing complete blocks, but the reader is trying to read complete blocks, the reader may keep getting partial blocks anyway.

The following example program demonstrates how to set up a one-way pipe between two related processes. Note that the processes MUST be related (parent, child, grandchild, etc.) since the pipe mechanism is based on the fact that file descriptors are inherited when a process is created. Error checking in the following program has been minimized in order to keep the code uncluttered and readable. In a "real" program more error checking on the system calls should be done.

```

/* who_wc.c */
/* demonstrates a one-way pipe between two processes.
   This program implements the equivalent of the shell command:

   who | wc -l

   which will count the number of users logged in.
*/

#include <stdio.h>

/* Define some manifest constants to make the code more understandable */

#define ERR    (-1)           /* indicates an error condition */
#define READ   0             /* read end of a pipe */
#define WRITE  1             /* write end of a pipe */
#define STDIN  0             /* file descriptor of standard in */
#define STDOUT 1             /* file descriptor of standard out */

int main()
{
    int    pid_1,             /* will be process id of first child - who */
          pid_2,             /* will be process id of second child - wc */
          pfd[2];            /* pipe file descriptor table. */

    if ( pipe ( pfd ) == ERR )           /* create a pipe */
    {                                     /* must do before a fork */
        perror ( " " );
        exit (ERR);
    }

    if (( pid_1 = fork ( ) ) == ERR)      /* create 1st child */
    {
        perror ( " " );
        exit (ERR);
    }

    if ( pid_1 != 0 )                    /* in parent */
    {
        if (( pid_2 = fork ( ) ) == ERR) /* create 2nd child */
        {
            perror ( " " );
            exit (ERR);
        }

        if ( pid_2 != 0 )                /* still in parent */
        {
            close ( pfd [READ] );         /* close pipe in parent */
            close ( pfd [WRITE] );        /* conserve file descriptors */
            wait ( ( int * ) 0 );         /* wait for children to die */
            wait ( ( int * ) 0 );
        }
    }
}

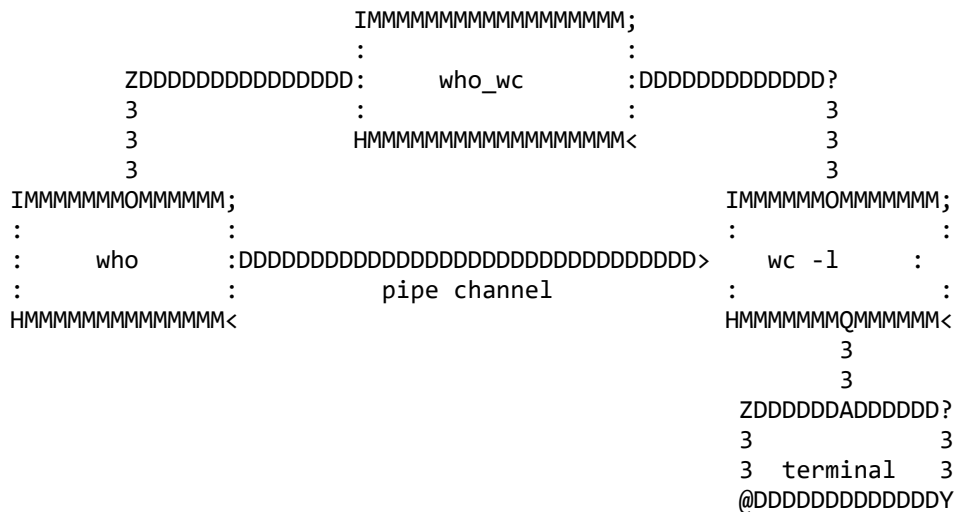
```

```

else                                     /* in 2nd child */
{
    close (STDIN);                      /* close standard input */
    dup ( pfd [READ] );                 /* read end of pipe becomes stdin */
    close ( pfd [READ] );               /* close unneeded I/O */
    close ( pfd [WRITE] );              /* close unneeded I/O */
    execl ("/bin/wc", "wc", "-l", (char *) NULL);
}
}
else                                     /* in 1st child */
{
    close (STDOUT);                    /* close standard out */
    dup ( pfd [WRITE] );                /* write end of pipes becomes stdout */
    close ( pfd [READ] );               /* close unneeded I/O */
    close ( pfd [WRITE] );              /* close unneeded I/O */
    execl ("/bin/who", "who", (char *) NULL);
}
exit (0);
}

```

The following is a diagram of the processes created by who_wc.



File Status

stat() - fstat()

The i-node data structure holds all the information about a file except the file's name and its contents. Sometimes your programs need to use the information in the i-node structure to do some job. You can access this information with the stat() and fstat() system calls. stat() and fstat() return the information in the i-node for the file named by a string and by a file descriptor, respectively. The format for the i-node struct returned by these system calls is defined in /usr/include/sys/stat.h. stat.h uses types built with the C language typedef construct and defined in the file /usr/include/sys/types.h, so it too must be included and must be included before the inclusion of the stat.h file.

The prototypes for stat() and fstat() are:

```
#include <sys/types.h>
#include <sys/stat.h>
```

```
int stat(file_name, stat_buf)
char *file_name;
struct stat *stat_buf;
```

```
int fstat(file_descriptor, stat_buf)
int file_descriptor;
struct stat *stat_buf;
```

where file_name names the file as an ASCII string and file_descriptor names the I/O channel and therefore the file. Both calls returns the file's specifics in stat_buf. stat() and fstat() fail if any of the following conditions hold:

- a path name component is not a directory (stat() only).

- file_name does not exist (stat() only).

- a path name component is off-limits (stat() only).

- file_descriptor does not identify an open I/O channel (fstat() only).

- stat_buf points to an invalid address.

Following is an extract of the stat.h file from the University's HP-9000. It shows the definition of the stat structure and some manifest constants used to access the st_mode field of the structure.

```

/* stat.h */

struct  stat
{
    dev_t      st_dev;          /* The device number containing the i-node */
    ino_t      st_ino;          /* The i-number */
    unsigned short st_mode;     /* The 16 bit mode */
    short      st_nlink;        /* The link count; 0 for pipes */
    ushort     st_uid;          /* The owner user-ID */
    ushort     st_gid;          /* The group-ID */
    dev_t      st_rdev;         /* For a special file, the device number */
    off_t      st_size;         /* The size of the file; 0 for special files */
    time_t     st_atime;        /* The access time. */
    int        st_spare1;
    time_t     st_mtime;        /* The modification time. */
    int        st_spare2;
    time_t     st_ctime;        /* The status-change time. */
    int        st_spare3;
    long       st_blksize;
    long       st_blocks;
    uint       st_remote:1;      /* Set if file is remote */
    dev_t      st_netdev;        /* ID of device containing */
    /* network special file */
    ino_t      st_netino;        /* Inode number of network special file */
    long       st_spare4[9];
};

#define S_IFMT 0170000 /* type of file */
#define S_IFDIR 0040000 /* directory */
#define S_IFCHR 0020000 /* character special */
#define S_IFBLK 0060000 /* block special */
#define S_IFREG 0100000 /* regular (ordinary) */
#define S_IFIFO 0010000 /* fifo */
#define S_IFNWK 0110000 /* network special */
#define S_IFLNK 0120000 /* symbolic link */
#define S_IFSOCK 0140000 /* socket */
#define S_ISUID 0004000 /* set user id on execution */
#define S_ISGID 0002000 /* set group id on execution */
#define S_ENFMT 0002000 /* enforced file locking (shared with S_ISGID)*/
#define S_ISVTX 0001000 /* save swapped text even after use */

```

Following is an example program demonstrating the use of the stat() system call to determine the status of a file:

```

/* status.c */
/* demonstrates the use of the stat() system call to determine the
   status of a file.
*/

#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>

#define ERR    (-1)
#define TRUE   1
#define FALSE  0

int main();

int main(argc, argv)
int argc;
char *argv[];
{
    int isdevice = FALSE;
    struct stat stat_buf;

    if (argc != 2)
    {
        printf("Usage:  %s filename\n", argv[0]);
        exit (1);
    }
    if ( stat( argv[1], &stat_buf) == ERR)
    {
        perror("stat");
        exit (1);
    }
    printf("\nFile:  %s  status:\n\n",argv[1]);
    if ((stat_buf.st_mode & S_IFMT) == S_IFDIR)
        printf("Directory\n");
    else if ((stat_buf.st_mode & S_IFMT) == S_IFBLK)
    {
        printf("Block special file\n");
        isdevice = TRUE;
    }
    else if ((stat_buf.st_mode & S_IFMT) == S_IFCHR)
    {
        printf("Character special file\n");
        isdevice = TRUE;
    }
    else if ((stat_buf.st_mode & S_IFMT) == S_IFREG)
        printf("Ordinary file\n");
    else if ((stat_buf.st_mode & S_IFMT) == S_IFIFO)
        printf("FIFO\n");
}

```



```

if (isdevice)
    printf("Device number:%d, %d\n", (stat_buf.st_rdev > 8) & 0377,
        stat_buf.st_rdev & 0377);
printf("Resides on device:%d, %d\n", (stat_buf.st_dev > 8) & 0377,
    stat_buf.st_dev & 0377);
printf("I-node: %d; Links: %d; Size: %ld\n", stat_buf.st_ino,
    stat_buf.st_nlink, stat_buf.st_size);
if ((stat_buf.st_mode & S_ISUID) == S_ISUID)
    printf("Set-user-ID\n");
if ((stat_buf.st_mode & S_ISGID) == S_ISGID)
    printf("Set-group-ID\n");
if ((stat_buf.st_mode & S_ISVTX) == S_ISVTX)
    printf("Sticky-bit set -- save swapped text after use\n");
printf("Permissions: %o\n", stat_buf.st_mode & 0777);

exit (0);
}

```

access()

To determine if a file is accessible to a program, the `access()` system call may be used. Unlike any other system call that deals with permissions, `access()` checks the real user-ID or group-ID, not the effective ones.

The prototype for the `access()` system call is:

```

int access(file_name, access_mode)
char *file_name;
int access_mode;

```

where `file_name` is the name of the file to which access permissions given in `access_mode` are to be applied. Access modes are often defined as manifest constants in `/usr/include/sys/file.h`. The available modes are:

Value	Meaning	file.h constant
----	-----	-----
00	existence	F_OK
01	execute	X_OK
02	write	W_OK
04	read	R_OK

These values may be ORed together to check for more than one access permission. The call to `access()` returns 0 if the program has the given access permissions, otherwise -1 is returned and `errno` is set to the reason for failure. This call is somewhat useful in that it makes checking for a specific permission easy. However, it only answers the question "do I have this permission?" It cannot answer the question "what permissions do I have?"

The following example program demonstrates the use of the `access()` system call to remove a file. Before removing the file, a check is made to make sure that the file exists and that it is writable (it will not remove a read-only file).

```
/* remove.c */

#include <stdio.h>
#include <sys/file.h>

#define ERR    (-1)

int main();

int main(argc, argv)
int argc;
char *argv[];
{
    if (argc != 2)
    {
        printf("Usage:  %s filename\n", argv[0]);
        exit (1);
    }
    if (access (argv[1], F_OK) == ERR)    /* check that file exists */
    {
        perror(argv[1]);
        exit (1);
    }
    if (access (argv[1], W_OK) == ERR)    /* check for write permission */
    {
        fprintf(stderr, "File:  %s  is write protected!\n", argv[1]);
        exit (1);
    }
    if (unlink (argv[1]) == ERR)
    {
        perror(argv[1]);
        exit (1);
    }
    exit (0);
}
```

Miscellaneous System Calls / Examples

Directories

A directory is simply a special file that contains (among other information) i-number/filename pairs. With the exception of 4.2 and 4.3 BSD, all versions of the UNIX system limit filenames to 14 characters. These short filenames make for a simple fixed size directory format on System V.

System V Directories

A directory contains structures of type `direct`, defined in the include file `/usr/include/sys/dir.h`. The include file `/usr/include/sys/types.h` must also be included to define the types used by the structure. The directory structure is:

```
#define DIRSIZ  14

struct direct {
    ino_t    d_ino;
    char     d_name[DIRSIZ];
};
```

It should be noted that the name of the file, `d_name` is NOT guaranteed to be null-terminated; programs should always be careful of this. Files which have been deleted will have i-numbers (`d_ino`) equal to zero; these should in general be skipped over when reading the directory. A directory is read simply by opening it (in read-only mode) and reading structures either one at a time or all at once. The following example program simply opens the current directory and prints the names of all the files it contains. The program simulates the `ls -a` command. Note that the file names are not sorted like the real `ls` command would do.

```
/* my_ls.c
   This program simulates the System V style ls -a command. Filenames
   are printed as they occur in the directory -- no sorting is done.
*/

#include <stdio.h>
#include <fcntl.h>
#include <sys/types.h>
#include <sys/dir.h>

#define ERR (-1)

int main()
{
    int fd;
    struct direct dir;

    if ((fd = open(".", O_RDONLY)) == ERR) /* open current directory */
    {
        perror("open");
        exit (1);
    }
}
```

```

while (( read (fd, &dir, sizeof (struct direct)) > 0 )
{
    if ( dir.d_ino == 0 )          /* is it a deleted file? */
        continue;                /* yes, so go read another */

    /* make sure we print no more than DIRSIZ characters */

    printf ("%.*s\n", DIRSIZ, dir.d_name);
}
close (fd);
exit (0);
}

```

If you need more information about the file such as size or permissions, you would use the stat() system call to obtain it.

Berkeley Style Directories

A directory contains structures of type direct, defined in the include file /usr/include/sys/ndir.h. The include file /usr/include/sys/types.h must also be included to define the types used by the structure. The directory structure is:

```

#define MAXNAMLEN 255
#define DIRSIZ_CONSTANT 14      /* equivalent to DIRSIZ */

struct direct {
    long    d_fileno;             /* file number of entry */
    short   d_reclen;             /* length of this record */
    short   d_namlen;             /* length of string in d_name */
    char    d_name[MAXNAMLEN + 1]; /* name (up to MAXNAMLEN + 1) */
};

#if !(defined KERNEL) && !(defined ATT3B2)
#define d_ino    d_fileno        /* compatibility */

```

Unlike on System V, filenames can be longer than 14 characters and the size of a directory structure can be variable. Therefore, the read() call can not be used to read the directory. Instead, Berkely style systems provide a set of library functions to read directories. These functions are also declared in the ndir.h include file. They are:

```

extern DIR *opendir();
extern struct direct *readdir();
extern long telldir();
extern void seekdir();
#define rewinddir(dirp) seekdir((dirp), (long)0)
extern void closedir();

```

The following example shows how to perform a Berkeley (or HP) style `ls -a` read of a directory. One important note: filenames in the directory structure are null-terminated in Berkeley style systems -- on System V they are not.

```
#include <stdio.h>
#include <sys/types.h>
#include <ndir.h>

main()
{
    DIR *dirp;
    struct direct *dp;

    dirp = opendir(".");           /* open the current directory */
    while ((dp = readdir(dirp)) != NULL)
    {
        if (dp->d_ino == 0)        /* ignore deleted files */
            continue;
        else
            printf("%s\n", dp->d_name); /* the name is null-terminated */
    }
}
```

For more information, type: `man directory` while logged onto the University's HP system.

Time

The UNIX operating system keeps track of the current date and time by storing the number of seconds that have elapsed since midnight January 1, 1970 UTC (Coordinated Universal Time, also known as Greenwich Mean Time (GMT)). This date is considered the informal "birthday" of the UNIX operating system. The time is stored in a signed long integer. (For the curious, assuming a 32 bit signed long integer, UNIX time will break at 03:14:08 January 19, 2038 UTC.)

In all versions of UNIX, the `time()` system call may be used to obtain the time of day. This call is peculiar in that if given the address of a long integer as an argument, it places the time in that integer and returns it. If, however, a null pointer is passed, the time of day is just returned.

Several routines are available to convert the long integer returned by `time()` into an ASCII date string. With the UNIX operating system, an ASCII date string is a string as shown below:

```
Day Mon dd hh:mm:ss yyyy
```

For example: Sat Mar 24 11:03:36 1990

The `ctime()` library function can be used to do the above conversion. An example is:

```
/* my_date.c
   print the current date and time in a format similar to the output
   of the date command.
*/

#include <stdio.h>
#include <time.h>      /* may need to be #include <sys/time.h> instead */

int main()
{
    long now, time();
    char *ctime();

    time (&now);
    printf("It is now %s\n", ctime (&now));

    exit (0);
}
```

Often you need access to specific information about the current date and time. The `localtime()` and `gmtime()` functions will provide it. They do this by converting the long integer returned by `time()` into a data structure called `tm`, which is defined in the `time.h` header file. In fact, this is what the header file looks like:

```
struct tm {
    int    tm_sec;    /* seconds after the minute - [0,59] */
    int    tm_min;    /* minutes after the hour - [0,59] */
    int    tm_hour;    /* hours since midnight - [0,23] */
    int    tm_mday;    /* day of the month - [1,31] */
    int    tm_mon;     /* months since January - [0,11] */
    int    tm_year;    /* years since 1900 */
    int    tm_wday;    /* days since Sunday - [0,6] */
    int    tm_yday;    /* days since January 1 - [0,365] */
    int    tm_isdst;   /* daylight savings time flag */
};
```

As you can see, there is quite a bit of information you can access. The `tm_isdst` member is non-zero if Daylight Savings Time is in effect. The `localtime()` function returns the time in the local time zone, whereas the `gmtime()` function returns the time in the UTC (or GMT) time zone. Both `localtime()` and `gmtime()` take as their argument a pointer to a long integer that represents the date and time as the number of seconds since January 1, 1970 (such a returned by `time()`). The return pointers to a `tm` structure, where the converted data is placed. The following example prints the local date in the familiar `mm/dd/yy` format:

```
/* day.c
   print date in mm/dd/yy format
*/

#include <stdio.h>
#include <time.h>      /* may need to be #include <sys/time.h> instead */

int main()
{
    long now, time();
    struct tm *today, *localtime();

    time (&now);
    today = localtime (&now);

    printf("Today is: %d/%d/%d\n", today->tm_mon + 1, today->tm_mday,
           today->tm_year);
    exit (0);
}
```

Parsing Input

When dealing with input from a command line, the first step is to parse (break up) the input line into tokens, which are groups of characters that form syntactic units; examples are words, strings, and special symbols. Following are some sample programs and functions that demonstrate various ways to parse an input line:

```
/* parse.c
    Split the input buffer into individual tokens.  Tokens are
    assumed to be separated by space or tab characters.
    A pointer to each token is stored in an array of pointers.
    This method is very similar to the argv argument to main().
*/
#include <stdio.h>
#define EVER ;;
#define MAXARG 64

int main()
{
    char buf[256];
    char *args[MAXARG];      /* accept MAXARG number of tokens */
    int num_arg,
        lcv;

    for (EVER)
    {
        printf("Enter line: ");
        if ((gets(buf)) == (char *) NULL)
        {
            putchar('\n');
            exit(0);
        }
        num_arg = parse_cmd(buf, args);
        printf("Number of tokens = %d\n", num_arg);
        for (lcv = 0; lcv < num_arg; lcv++)
            puts(args[lcv]);
    }
}
```



```

int parse_cmd(buf, args)
char *buf;
char **args;
{
    int count = 0;

    while (*buf != '\0' && count < MAXARG)
    {
        while ((*buf == ' ') || (*buf == '\t'))
            *buf++ = '\0';
        *args++ = buf;
        ++count;
        while ((*buf != '\0') && (*buf != ' ') && (*buf != '\t'))
            buf++;
    }
    *args = (char *) NULL; /* make the last element of the array null */
    return(count);        /* return the number of tokens parsed */
}

```

There is a C library function available that makes parsing a string into tokens very easy; it is `strtok()`. Following is the above function (`parse_cmd()`) as implemented using `strtok()`:

```

int parse_cmd(line, args)
char *line;
char *args[];
{
    int count = 0;
    char *str, *strtok();
    static char delimiters[] = " \t\n";

    while ((str = strtok(line, delimiters)) != (char *) NULL)
    {
        line = (char *) NULL;
        args[count++] = str;
    }
    args[count] = (char *) NULL;
    return(count);
}

```

`strtok()` takes as arguments a pointer to the input string and a pointer to a string containing the character or characters that delimit the token. In the above example, the delimiters were defined to be a space, tab, or newline. You are free to change the delimiter at any time. If you wish `strtok()` to parse the complete line, you must pass a null-pointer on the second and subsequent calls to `strtok()` (note that "line" was set to null inside the body of the while loop). `strtok()` returns a null-pointer when the end of the input string is reached.

`strtok()` is very useful in parsing the individual path elements as defined in the PATH environment variable (set the delimiter to ":").

CURSES

What is curses? curses is a terminal-independent library of C routines and macros that you use to write "window-based" screen management programs on the UNIX system. curses is designed to let programmers control terminal I/O in an easy fashion. Providing an easy-to-use "human interface" for users is an increasingly important requirement for operating systems. Such a connection between the machine and the humans that use it plays an important role in the overall productivity of the system. curses gets its name from what it does: cursor manipulation.

What can curses do? Among the functions to be found in curses are those that:

- Move the cursor to any point on the screen
- Insert text anywhere on the screen, doing it even in highlight mode
- Divide the screen into rectangular areas called windows
- Manage each window independently, so you can be scrolling one window while erasing a line in another
- Draw a box around a window using a character of your choice
- Write output to and read input from a terminal screen
- Control the data output and input -- for example, to print output in bold type or prevent it from echoing (printing back on a screen)
- Draw simple graphics

If these features leave you unimpressed, remember that they are only tools. When you use these tools in your programs, the results can be spectacular. The point is -- curses is easy to use and ready to go -- so that you can concentrate on what you want your program to do. curses will make you program look sharp.

Where did curses come from? The author of curses is Ken Arnold who wrote the package while a student at the University of California, Berkeley. At the same time, Bill Joy was writing his editor program, vi. Ken Arnold credits Bill Joy with providing the ideas (as well as code) for creating the capability to generally describe terminals, writing routines to read the terminal database, and implementing routines for optimal cursor movement. The original source of information about curses is Ken Arnold's paper entitled "Screen Updating and Cursor Movement Optimization: A Library Package".

What makes curses tick? The original version of curses developed by Ken Arnold incorporated a database known as termcap, or the terminal capabilities database. In System V Release 2, the termcap database was replaced by the terminfo data base, and curses was rewritten to incorporate it. Both of these versions of curses can be used with more than one hundred terminals. The information in the termcap or terminfo database is used by the curses routines to determine what sequence of special characters must be sent to a particular terminal to cause it to clear the screen, move the cursor up one line, delete a line, etc. It is these databases that make curses truly terminal independent, since any terminal not already in the database can be added by a system administrator, and since the structure of both databases allows users to add their own local additions or modifications for a particular terminal.

How to use curses -- the basics: There are a couple of things you have to know before you can start using the curses library. First, when you compile a C program that call curses routines, you must specify to the cc command that the curses library is to be linked in with the program. This is done with the -lcurses option, which must be specified after all the C program files. The following is an example cc command line for use on systems that support the terminfo database:

```
cc myprog.c -lcurses
```

Next is an example cc command line for use on systems that support the termcap database:

```
cc myprog.c -lcurses -ltermcap
```

Second, all program files that reference curses routines must include the header file <curses.h> <curses.h> will include the header <stdio.h> so it is not necessary for your program to include it. It won't hurt anything if you do -- it just slows down the compilation. Lastly, before you run a program that uses curses, you must inform curses what type of terminal you have. You do this by setting the shell variable TERM to the type of terminal you are using (e.g. a DEC VT100) and exporting the TERM variable into the environment. This is done in the following manner:

```
$ TERM=vt100
$ export TERM
```

This action is usually done for you by your .profile when you log it.

The <curses.h> header file contains declarations for variables, constants, data structures and macros. Among the variables are two integer variables that prove to be very useful: LINES and COLS. LINES is automatically set to the number of lines on your terminal; COLS is set to the number of columns. Many of the curses routines address the terminal's screen, in that they move the cursor to a specific place, or address. This address is specified as a particular row and column (specified as arguments to the routine), where the address of the upper left-hand corner is row LINES-1 and column COLS-1 (LINES-1, COLS-1). Following is a layout of the terminal screen:

```

IMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMM;
:(0,0)                                     (0, COLS-1):
:                                           :
:                                           :
:                                           :
R :                                           :
O :                                           :
W :                                           :
:                                           :
:               .               (row,col)      :
:                                           :
:                                           :
:                                           :
:                                           :
:(LINES-1,0)                             (LINES-1, COLS-1):
HMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMMM<
                                COLUMN

```

All programs using the curses library must have the following basic structure:

```
#include <curses.h>

main()
{
    initscr();

    /* main program */

    endwin();
}
```

The `initscr()` function must be called before any other curses routines. Its function is to determine the terminal type from the `TERM` environment variable and to initialize certain data structures and variables (e.g. `LINES` and `COLS`). The `endwin()` function should be called prior to program exit to restore the terminal's original state. Some curses routines change the terminal's characteristics (e.g. go into raw mode and turn off echoing) and must be undone before the program exits; otherwise the terminal is left in an odd state and the user may not know how to change it back. Here is how to fix the terminal if a curses program leaves it in a "funny" state:

on System III and System V UNIX systems (including XENIX) type:

```
stty sane ctrl-j
```

note that you must type a control-j and not the return key, since most likely `NEWLINE` mapping will be off and the `RETURN` key will not work.

on Berkeley UNIX systems type:

```
stty -raw -cbreak -nl echo ek ctrl-j
```

Windows and screens: Conceptually, a window is an independent rectangular area of characters displayed on the screen. Physically, a window is a `WINDOW`, that is, a C data structure that holds all the information about a window.

The Standard Screen - `stdscr`: The traditional definition of the "standard screen" is a window or a set of windows that fills the entire screen of a video display terminal. The structure that describes `stdscr` is a `WINDOW`, or more precisely, a pointer to a `WINDOW`. A `WINDOW` is a character array that maintains an image of the terminal screen, known as the screen image. The screen image array in `stdscr` is automatically made the length and width of the terminal screen. Thus, there is one character in that array for every place on the screen.

The Current Screen - `curscr`: `curses` does not know directly what the terminal is displaying; it would be even slower to have to query the terminal to find out what character is being displayed at each location. Instead, `curses` keeps an image of what it thinks the screen looks like in a window called `curscr`. `curscr`, like `stdscr`, is created automatically when you initialize `curses` with `initscr()`. `curscr` is a `WINDOW`, and has a screen image the size of the physical screen. When `refresh()` is called, it writes the characters that it is sending to the terminal into their corresponding location in the screen image of `curscr`. `curscr` contains the image of the screen as `curses` thinks it was made to look by the last `refresh()`. `refresh()` uses the screen image in `curscr` to minimize its work. When it goes to refresh a window, it compares the contents of that window to `curscr`. `refresh()` assumes that the physical screen looks like `curscr` so it does not output characters that are the same in `curscr` and the window that is being refreshed. In this way, `refresh()` minimizes the number of characters that it sends to the screen and save a great deal of time.

The following are a few of the more commonly used `curses` routines. The list is not comprehensive:

Terminal Modes: the terminal modes for I/O are usually set after the call to `initscr()`. None of the mode setting routines accept parameters.

<code>echo()</code> / <code>noecho()</code>	These functions allow programmers to turn on or off the terminal driver's echoing to the terminal. The default state is echo on. The function <code>noecho()</code> disables the automatic echoing.
<code>nl()</code> / <code>nonl()</code>	These functions allow programmers to enable or disable carriage return/newline mappings. When enabled, carriage return is mapped on input to newline and newline is mapped on output to newline/carriage return. The default state is mapping enabled., and <code>nonl()</code> is used to turn this mapping off. It is interesting to note that while mapping is disabled, cursor movement is optimized.
<code>cbreak()</code> / <code>nocbreak()</code>	Canonical processing (line at a time character processing) is disabled within the terminal driver when calling <code>cbreak()</code> , allowing a break for each character. Interrupt and flow control keys are unaffected. The default state is <code>nocbreak</code> , which enables canonical processing.
<code>raw()</code> / <code>noraw()</code>	These functions are similar to the <code>cbreak()</code> / <code>nocbreak()</code> functions, except that interrupt and flow control key are also disabled or enabled.
<code>savetty()</code> / <code>resetty()</code>	The current state of the terminal can be saved into a buffer reserved by <code>curses</code> when calling <code>savetty()</code> function. The last save state can be restored via the <code>resetty()</code> function.
<code>gettmode()</code>	This function is used to establish the current tty mode while in <code>curses</code> . It reads the baud rate of the terminal, turns off the mapping of carriage returns to line feeds on output, and the expansion of tabs into spaces by the system.

I/O Function:

addch() This function adds a character to a window at the current cursor position.

```
#include <curses.h>
main()
{
    initscr();
    addch('e');
    refresh();
    endwin();
}
```

mvaddch() This function moves a character into a window at the position specified by the x and y coordinates.

```
#include <curses.h>
main()
{
    int x,y;

    x = 3; y = 10;
    initscr();
    mvaddch(x, y, 'e');
    refresh();
    endwin();
}
```

addstr() This function adds the specified string to a window at the current cursor position.

```
#include <curses.h>
main()
{
    initscr();
    addstr("This is a string example.");
    refresh();
    endwin();
}
```

mvaddstr() This function moves the specified string into a window located at the position specified by the x and y coordinates.

```
#include <curses.h>
main()
{
    int x,y;

    x = 3; y = 10;
    initscr();
    mvaddstr(x, y, "This is the string example.");
    refresh();
    endwin();
}
```

`printw()` This function outputs formatted strings at the current cursor position and is similar to the `printf()` function of C, in that multiple arguments may be specified.

```
#include <curses.h>
main()
{
    static char *word = "example";
    int number = 1;
    initscr();
    printw("this is just %d %s of a formatted string!\n",number,word);
    refresh();
    endwin();
}
```

`mvprintw()` This function outputs formatted strings at the line specified in `y` and the column specified in `x`. Multiple arguments may be given.

```
#include <curses.h>
main()
{
    static char *word = "example";
    int number = 1;
    int x = 3, y = 10;
    initscr();
    mvprintw(x ,y, "this is just %d %s of a formatted string!\n",
              number,word);
    refresh();
    endwin();
}
```

`move()` This function moves the cursor to the line/column coordinates given.

```
#include <curses.h>
main()
{
    int line = 3, column = 10;
    initscr();
    move(line, column);
    refresh();
    endwin();
}
```

`getyx()` This function is used to determine and return the current line/column location of the cursor.

```
#include <curses.h>
main()
{
    WINDOW *win;
    int y, x;
    initscr();
    win = newwin(10,5,12,39);
    getyx(win, y, x)
    refresh();
    endwin();
}
```

`getch()` This function is used to read a single character from the keyboard, and returns an integer value. It is similar to the the C standard I/O function `getc()`;

```
#include < curses.h>
main()
{
    int in_char;
    initscr();
    in_char = getch();
    refresh();
    endwin();
}
```

`inch()` This function returns the character from under the current cursor position of the terminals screen, in an integer.

```
#include < curses.h>
main()
{
    int in_char;
    initscr();
    in_char = inch();
    refresh();
    endwin();
}
```

`mvinch()` This function is used to get the character under the cursor location specified as x and y coordinates. The value returned is an integer.

```
#include < curses.h>
main()
{
    int in_char;
    initscr();
    in_char = mvinch(3, 10);
    refresh();
    endwin();
}
```

`clear()` This function completely clear the terminal screen by writing blank spaces to all physical screen locations via calls to `erase()` and `clearok()`, and is completed by the next call to `refresh()`.

```
#include < curses.h>
main()
{
    initscr();
    clear();
    refresh();
    endwin();
}
```


`erase()` This function is used to insert blank spaces in the physical screen and, like `clear()`, erases all data on the terminal screen, but does not require a call to `refresh()`.

```
#include <curses.h>
main()
{
    initscr();
    erase();
    endwin();
}
```

`clrtobot()` This function is used to clear the physical screen from the current cursor position to the bottom of the screen, filling it with blank spaces.

```
#include <curses.h>
main()
{
    initscr();
    clrtobot();
    refresh();
    endwin();
}
```

`clrtoeol()` This function is used to clear the physical screen from the current cursor position to the end of the physical screen line by filling it with blank spaces.

```
#include <curses.h>
main()
{
    initscr();
    clrtoeol();
    refresh();
    endwin();
}
```

`delch()` This function deletes the character under the current cursor position, moving all characters on that line (located to the right of the deleted character) one position to the left, and fills the last character position (on that line) with a blank space. The current cursor position remains unchanged.

```
#include <curses.h>
main()
{
    initscr();
    delch();
    refresh();
    endwin();
}
```

`mvdelch()` This function deletes the character under the cursor position at the line/column specified in `y/x`. In all other respects, it works the same as the `delch()` function,

```
#include <curses.h>
main()
{
    initscr();
    mvdelch(3, 10);
    refresh();
    endwin();
}
```

`insch()` This function is used to insert the character named in `'c'` to be inserted at the current cursor position, causing all characters to the right of the cursor (on that line, only) to shift one space to the right, losing the last character of that line. The cursor is moved one position to the right of the inserted character.

```
#include <curses.h>
main()
{
    initscr();
    insch('c');
    refresh();
    endwin();
}
```

`mvinsch()` This function inserts the character named in `'c'` to the line/column position named in `y/x`, and otherwise works identically to the `insch()` function.

```
#include <curses.h>
main()
{
    initscr();
    mvinsch(3, 10, 'c');
    refresh();
    endwin();
}
```

`deleteln()` This function allows the deletion of the current cursor line, moving all lines located below up one line and filling the last line with blank spaces. The cursor position remains unchanged.

```
#include <curses.h>
main()
{
    initscr();
    deleteln();
    refresh();
    endwin();
}
```

`insertln()` This function inserts a blank filled line at the current cursor line, moving all lines located below down one line. The bottom line is lost, and the current cursor position is unaffected.

```
#include <curses.h>
main()
{
    initscr();
    insertln();
    refresh();
    endwin();
}
```

`refresh()` This function is used to update the physical terminal screen from the window buffer and all changes made to that buffer (via curses functions) will be written. If the buffer size is smaller than the physical screen, then only that part of the screen is refreshed, leaving everything else unchanged.

```
#include <curses.h>
main()
{
    initscr();
    /* curses function call(s) here */
    refresh();
    endwin();
}
```

`wrefresh()` This function is identical to the `refresh()` function, except that the refresh operation is performed on the named window.

```
#include <curses.h>
main()
{
    WINDOW *win;

    initscr();
    /* curses function call(s) here */
    wrefresh(win);
    endwin();
}
```

`initscr()` This function call must be present in all programs calling the curses functions. It clears the physical terminal screen and sets up the default modes. It should be the first call to the curses functions when using the library to initialize the terminal.

`endwin()` This function call should be present in any program using the curses functions, and should also be the last function call of that program. It restores all terminal settings to their original state prior to using the `initscr()` function call and it places the cursor to the lower left hand portion of the screen and terminates a curses program.

`attrset()` This function allows the programmer to set single or multiple terminal attributes. The call `attrset(0)` resets all attributes to their default state.

```
#include <curses.h>
main()
{
    initscr();
    attrset(A_BOLD);
    /* sets character attributes to bold */
    ...
    /* curses function call(s) here */
    attrset(0);
    /* resets all attributes to default */
    refresh();
    endwin();
}
```

`attron()` This function is used to set the named attribute of a terminal to an on state.

```
#include <curses.h>
main()
{
    initscr();
    attron(A_BOLD);
    /* sets character attribute to bold */
    ...
    /* curses function call(s) here */
    refresh();
    endwin();
}
```

`attroff()` This function is the opposite of the `attron()` function and will turn off the named attribute of a terminal.

```
#include <curses.h>
main()
{
    initscr();
    attron(A_BOLD);
    /* sets character attribute to bold */
    ...
    /* curses function call(s) here */
    attroff(A_BOLD);
    /* turns off the bold character attribute */
    refresh();
    endwin();
}
```

`standout()` This function sets the attribute `A_STANDOUT` to an on state, and is nothing more than a convenient way of saying `attron(A_STANDOUT)`.

```
#include <curses.h>
main()
{
    initscr();
    standout();
    ...
    /* curses function call(s) here */
    refresh();
    endwin();
}
```

`standend()` This function, like `standout()`, is just a convenient way of saying `attroff(A_STANDOUT)`, meaning that the `A_STANDOUT` attribute is set to an off state. Actually, this function resets all attributes to the off state.

```
#include <curses.h>
main()
{
    initscr();
    standout();
    ...
    /* curses function call(s) here */
    standend();
    /* end of attribute settings */
    refresh();
    endwin();
}
```

`box()` This function draws a box around the edge of the window. One of its arguments is the horizontal character and the other is the vertical character.

```
#include <curses.h>
main()
{
    initscr();
    box(stdscr, '-', '*');
    /* draws a box around the stdscr */
    /* horizontal characters are '-' and vertical characters are '*' */
    refresh();
    endwin();
}
```

Attribute Values: the following is a list of the terminal attributes that may be set on or off using the curses library. It is important to note that all of these attributes may not be available to the physical terminal, depending upon the given terminal's characteristics.

A_STANDOUT - this attribute allows the terminal to display characters in highlight, bold, or some other fashion (depending upon the terminal's characteristics).

A_REVERSE - this attribute allows the terminal to display its characters in reverse video.

A_BOLD - this attribute allows the terminal to display its characters in bold lettering.

A_DIM - this attribute allows the terminal to display its characters at less intensity than normal.

A_UNDERLINE - this attribute allows the terminal to display characters with a horizontal line beneath them (underlined).

A_BLINK - this attribute allows the terminal to display blinking characters that will appear and disappear at a rate depending upon the terminal characteristics.

Creating and Removing Windows:

WINDOW *newwin(lines, cols, y1, x1) will create a new window. The new window will have lines lines and cols columns, with the upper left corner located at (y1,x1). newwin() returns a pointer to WINDOW that points at the new window structure. The screen image in the new window is filled with blanks.

WINDOW *subwin(win, lines, cols, y1, x1) will create a sub-window. win is a pointer to the parent window. The other arguments are the same as in newwin(), except lines and cols are interpreted relative to the parent's window and not the terminal screen. A sub-window is a real WINDOW and may have sub-windows just as easily as the parent window.

delwin(win) will delete the specified window. delwin() calls the system utility free() to return the space to the pool of available memory. If the window is a sub-window, delwin() does not free() the space because that space is still being used by the parent. Deletint a parent does not free the space occupied by sub-windows. The sub-windows will continue to occupy space, but their screen images will be undefined. You should take care to delete windows in the proper order and when needed in order to maintain good housekeeping of the available memory.

Window Specific Functions: these functions are some of the functions above applied to a window. A 'w' is placed before the function name, and the first argument is a pointer to the window.

waddch(win, ch)	winch(win)
waddstr(win, str)	winsch(win, c)
wclear(win)	winsertln(win)
wclrtoobot(win)	wmove(win, y, x)
wclrtoeol(win)	wprintw(win, fmt, arg1, arg2, ...)
wdelch(win, c)	wrefresh(win)
wdeleteln(win)	wscanw(win, fmt, agr1, arg2, ...)
werase(win)	wstandout(win)
wgetch(win)	wstandend(win)
wgetstr(win, str)	

Move and Act Function: these functions first move the cursor, then perform their action. The function names have a 'mv' placed before the corresponding function above.

mvaddch(y, x, ch)	mvwaddch(win, y, x, ch)
mvaddstr(y, x, str)	mvwaddstr(win, y, x, ch)
mvdelch(y, x)	mvwdelch(win, y, x)
mvdeleteln(y, x)	mvwdeleteln(win, y, x)
mvinch(y, x)	mvwinch(win, y, x)
mvinsch(y, x, ch)	mvwinsch(win, y, x, ch)
mvinsertln(y, x)	mvwininsertln(win, y, x)

The following example program demonstrates a few of the curses functions:

```
/* disptime.c
   this program displays the time and refreshes the screen once
   every second, so that the screen resembles a digital clock.
*/

#include <curses.h>
#include <time.h>
#include <signal.h>
#define EVER ;;

main()
{
    void sig_catch();
    long seconds;
    static char *title = "The current time is", *convtime, *ctime();

    /* call sig_catch if the user hits DELETE/BREAK key */
    signal (SIGINT, sig_catch);
    /* initial setup of curses */
    initscr();

    /* output title centered */
    mvaddstr (LINES / 2-1, (COLS - strlen (title)) / 2, title);

    for (EVER)
    {
        /* get time and convert to ASCII */
        time (&seconds);
        convtime = ctime (&seconds);
        /* display time centered under the title */
        mvaddstr (LINES / 2 , (COLS - strlen (convtime)) / 2, convtime);
        refresh ();
        sleep (1);
    }

    /* signal handling routine, call endwin() and exit */

    void sig_catch()
    {
        endwin ();
        exit (1);
    }
}
```


AWK

What is awk?

awk is one of the more unusual UNIX commands. Named after and by its creators: Aho, Weinberger, and Kernighan, awk combines pattern matching, comparison making, line decomposition, numerical operations, and C-like programming features into one program.

awk is a "small" language, in that it lacks some of the more complicated features found in traditional languages like C, Pascal, and Ada. In general, awk omits many mechanisms that support the development of large applications, such as modules and user defined types.

Nonetheless, awk is a powerful and general-purpose language, capable of nearly anything you would want a programming language to do. The language omissions foster the development of small applications, as do the robust string manipulation capabilities and the powerful table facility. awk's automatic storage management frees the programmer from having to explicitly keep track of memory -- this alone can cut programming and debugging in half.

The Structure of an awk Program

Many applications consist of simple collections of patterns and actions. Each time a pattern is recognized in the input, the corresponding action is executed. C code to do this would resemble the following:

```
while ( getRecord() != EOF) {  
    if ( pattern1) { action1 }  
    if ( pattern2) { action2 }  
    ...  
    if ( patternN) { actionN }  
}
```

Because this approach is suitable to so many applications, awk's syntax was specifically streamlined to support it. The corresponding awk program eliminates the outermost control-flow syntax:

```
pattern1 { action1 }  
pattern2 { action2 }  
...  
patternN { actionN }
```

awk read each line in the input file(s) one at a time. When a line is read, each pattern is tested in sequence. Whenever a pattern matches the current line, the corresponding action is executed. This continues until all the input has been processed.

awk breaks each input record into fields separated by whitespace (or a specified delimiter). Within the awk program, fields are designated \$1, \$2, etc., and the variable NF is set to the total number of fields in the current record. The variable \$0 stands for the entire record not broken into fields.

Patterns resemble boolean expressions, with the addition of regular expression operators, ~, and a syntax for regular expressions contained in slashes added e.g. the pattern \$2 ~ /foo/ is true if the field \$2 contains a substring "foo". A regular expression without an explicit range is matched against the input record (\$0).

Several patterns are special. Action connected to BEGIN is executed once before any records are read. END action is executed once after all input records are processed. Omitted patterns match every record.

Example Application - Create an Index

An index might be appended to the end of a report, or it might be used to extract specific cross-reference information from source code. An awk script to generate an index is characteristically simple:

```
BEGIN { while (getline < "keywords" > 0)
        KEY[$0] = ""
        }
        { for (k in KEY)
            if( $0 ~ k)
                KEY[k] = KEY[k] " " NR
        }
END    { for ( k in KEY)
        print k, KEY[k]
        }
```

The action associated with BEGIN reads a list of keywords to be indexed from a file called keywords. Each keyword is used as an index into a table called KEY, and the corresponding entries in the table are nulled.

The middle action (with no pattern) executes for every input file line. Each line is checked for the presence of each keyword in the KEY table. The line number (NR) is appended to the KEY entry for each match.

After all lines are processed, the END action prints each keyword followed by the lines where it appeared.

Rapid Prototyping

awk is suitable for prototyping applications -- quickly implementing ideas to test feasibility before making a major investment in implementation. If the idea is a bad one, this can be discovered after 50 lines of awk rather than 5000 lines of C. Rapid prototyping provides prospective users with a program to "play with" -- the most effective way to find out what the users really want in the final product. Typical awk prototypes are often less than 10% of the length of the equivalent C program. Once you have a suitable prototype program written in awk, you can decide whether to recode the program in a conventional language like C, or just stick with the awk program itself.

References

1. UNIX System V Programmer's Guide
Prentice-Hall, 1987
2. Using C on the UNIX System
David C. Curry
O'Reilly & Associates, Inc. 1989
3. UNIX for MS-DOS Programmers
Steven Mikes
Addison-Wesley, 1989
4. Advanced Programmer's Guide to UNIX System V
Rebecca Thomas, Lawrence Rogers, Jean Yates
Osborne/McGraw-Hill, 1986
5. The Design of the UNIX Operating System
Maurice Bach
Prentice-Hall, 1986
6. Topics in C Programming
Stephen Kochan and Patrick Wood
Hayden Books, 1987
7. Advanced UNIX Programming
Marc Rochkind
Prentice-Hall, 1985
8. Programming with Curses
John Strang
O'Reilly & Associates, 1986